

Variational Treatment of Probabilistic Directed Graphical Models

Satish Kumar

July 27, 2018

LEAP Lab, IISC Bengaluru

Table of contents

1. Generative modeling
2. Evidence Lower Bound(ELBO)
3. Variational Inference in Deep Learning
4. Variational Autoencoder
5. Reparameterization trick

Generative modeling

Density Estimation



Figure 1: Some generative models perform density estimation. These models take a training set of examples drawn from an unknown data-generating distribution p_{data} and return an estimate of that distribution. The estimate p_{model} can be evaluated for a particular value of \mathbf{x} to obtain an estimate $p_{\text{model}}(\mathbf{x})$ of the true density $p_{\text{model}}(\mathbf{x})$. This figure illustrates the process for a collection of samples of one-dimensional data and a Gaussian model.

An Ideal Generative model



Training examples

Model samples

Why study Generative modeling?

- Training and sampling from generative models is an excellent test of our ability to represent and manipulate high-dimensional probability distributions.
- Generative models can be trained with missing data and can provide predictions on inputs that are missing data.
- It enable machine learning to work with multi-modal outputs
- Realistic generation of samples from some distribution.

Evidence Lower Bound(ELBO)

EM:A Latent Variable View

Consider a probabilistic model: Observed variables \mathbf{X} and latent variables \mathbf{Z} . Our goal is to minimize the likelihood function given by

$$p(\mathbf{X}|\boldsymbol{\theta}) = \sum_{\mathbf{Z}} p(\mathbf{X}, \mathbf{Z}|\boldsymbol{\theta})$$

ELBO: Evidence Lower Bound

- We introduce a distribution $q(\mathbf{z})$ defined over latent variables, and for any choice of $q(\mathbf{z})$ the following decomposition holds

$$\ln p(\mathbf{X}|\boldsymbol{\theta}) = \mathcal{L}(q, \boldsymbol{\theta}) + \text{KL}(q\|p)$$


where we have defined

$$\begin{aligned}\mathcal{L}(q, \boldsymbol{\theta}) &= \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln \left\{ \frac{p(\mathbf{X}, \mathbf{Z}|\boldsymbol{\theta})}{q(\mathbf{Z})} \right\} \\ \text{KL}(q\|p) &= - \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln \left\{ \frac{p(\mathbf{Z}|\mathbf{X}, \boldsymbol{\theta})}{q(\mathbf{Z})} \right\}.\end{aligned}$$

- using the product rule of probability

$$\ln p(\mathbf{X}, \mathbf{Z}|\boldsymbol{\theta}) = \ln p(\mathbf{Z}|\mathbf{X}, \boldsymbol{\theta}) + \ln p(\mathbf{X}|\boldsymbol{\theta})$$


Illustration of ELBO



E-step


Illustration of the E step of the EM algorithm. The q distribution is set equal to the posterior distribution for the current parameter values θ^{old} , causing the lower bound to move up to the same value as the log likelihood function, with the KL divergence vanishing.

$$\text{KL}(q||p) = 0$$




M-step

Illustration of the M step of the EM algorithm. The distribution $q(\mathbf{Z})$ is held fixed and the lower bound $\mathcal{L}(q, \theta)$ is maximized with respect to the parameter vector θ to give a revised value θ^{new} . Because the KL divergence is nonnegative, this causes the log likelihood $\ln p(\mathbf{X}|\theta)$ to increase by at least as much as the lower bound does.



Animation of Lower Bound

The EM algorithm involves alternately computing a lower bound on the log likelihood for the current parameter values and then maximizing this bound to obtain the new parameter values. See the text for a full discussion.



Variational Inference in Deep Learning

Some preliminaries

- We assume the observed variable \mathbf{x} is a random sample from an unknown underlying process where the distribution $p^*(\mathbf{x})$ is unknown
- we approximate this process with a chosen model $p_\theta(\mathbf{x})$ with parameter θ

$$\mathbf{x} \sim p_\theta(\mathbf{x})$$

- Learning is the process of searching the θ such that, for any observed variable \mathbf{x}

$$p_\theta(\mathbf{x}) \approx p^*(\mathbf{x})$$

- we wish $p_\theta(\mathbf{x})$ to be sufficiently flexible to be able to adapt to the data
- Often, such as in case of classification or regression problems, we are not interested in learning an unconditional model $p_\theta(\mathbf{x})$, but a conditional model $p_\theta(\mathbf{y}|\mathbf{x})$
- that approximates the underlying conditional distribution $p^*(\mathbf{y}|\mathbf{x})$

$$p_\theta(\mathbf{y}|\mathbf{x}) \approx p^*(\mathbf{y}|\mathbf{x})$$

Parameterizing Conditional distributions with Neural Nets

- We parameterize conditional distributions with neural networks
- In image classification, neural networks parameterize a categorical distribution $p_{\theta}(\mathbf{y}|\mathbf{x})$ over a class label \mathbf{y} , conditioned on an image \mathbf{x} .

$$\begin{aligned}\mathbf{p} &= \text{NeuralNet}(\mathbf{x}) \\ p_{\theta}(\mathbf{y}|\mathbf{x}) &= \text{Categorical}(\mathbf{y}; \mathbf{p})\end{aligned}$$

- where for all $p_i \in \mathbf{p}$
- and the last operation of the neural net is a *softmax*() function such that $\sum_i p_i = 1$

In the case of a Gaussian MLP as encoder or decoder, we let the encoder or decoder be a multivariate Gaussian with a diagonal covariance structure:

$$\log p(x|z) = \log \mathcal{N}(x; \mu, \sigma^2 \mathbf{I})$$

$$\text{where } \mu = W_4 \mathbf{h} + \mathbf{b}_4$$

$$\log \sigma^2 = W_5 \mathbf{h} + \mathbf{b}_5$$

$$\mathbf{h} = \tanh(W_3 \mathbf{z} + \mathbf{b}_3)$$

where $\{W_3, W_4, W_5, \mathbf{b}_3, \mathbf{b}_4, \mathbf{b}_5\}$ are the weights and biases of the MLP and part of θ when used as decoder. Note that when this network is used as an encoder $q_\phi(z|x)$, then z and x are swapped, and the weights and biases are variational parameters ϕ .

Learning in Fully Observed models with neural nets

- **Dataset:** We often collect a dataset \mathcal{D} consisting of $N \geq 1$ datapoints:

$$\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\} \equiv \{\mathbf{x}^{(i)}\}_{i=1}^N \equiv \mathbf{x}^{(1:N)}$$

- Under the i.i.d. assumption, the probability of the datapoints given the parameters factorizes as a product of individual datapoint probabilities.
- **Maximum Likelihood and Minibatch SGD**

$$\log p_{\theta}(\mathcal{D}) = \sum_{\mathbf{x} \in \mathcal{D}} \log p_{\theta}(\mathbf{x})$$

- Using automatic differentiation tools, we can efficiently compute gradients of this objective, and use such gradient to find the local optimum of the ML objective.
- We can opt either Stochastic Gradient or Batch Gradient Methods

Learning and Inference in Deep latent variable models

- **Latent Variables, \mathbf{z}** are variables that are part of the model, but which we don't observe, and are therefore **not part of the dataset**.
- In case of unconditional modeling of observed variable \mathbf{x} , the directed graphical model would then represent a joint distribution $p_{\theta}(\mathbf{x}, \mathbf{z})$ over both the observed variables \mathbf{x} and the latent variables \mathbf{z} .
- The marginal distribution over the observed variables $p_{\theta}(\mathbf{x})$, is given by:

$$p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}, \mathbf{z}) d\mathbf{z}$$

- This is also called the (single datapoint) **marginal likelihood** or the **model evidence**, when taking as a function of θ .

Deep Latent Variable Models(DLVMs)

- We use the term deep latent variable model (DLVM) to denote a latent variable model $p_{\theta}(\mathbf{x}, \mathbf{z})$ whose distributions are parameterized by neural networks
- Such a model can be conditioned on some context, like $p_{\theta}(\mathbf{x}, \mathbf{z} | \mathbf{y})$
- The simplest, and most common, graphical model with latent variables is one that is specified as factorization with the following structure:

$$p_{\theta}(\mathbf{x}, \mathbf{z}) = p_{\theta}(\mathbf{z}) p_{\theta}(\mathbf{x} | \mathbf{z})$$

- The distribution $p(\mathbf{z})$ is often called the prior distribution over \mathbf{z}

Example: DLVM for multivariate Bernoulli data

- For binary data \mathbf{x} , with a spherical Gaussian latent space, and a factorized Bernoulli observation model:

$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$

$$\mathbf{p} = \text{DecoderNeuralNet}_{\theta}(\mathbf{z})$$

$$\begin{aligned}\log p(\mathbf{x}|\mathbf{z}) &= \sum_{j=1}^D \log p(x_j|\mathbf{z}) = \sum_{j=1}^D \log \text{Bernoulli}(x_j; p_j) \\ &= \sum_{j=1}^D x_j \log p_j + (1 - x_j) \log(1 - p_j)\end{aligned}$$

- where $\forall p_j \quad \mathbf{p} : 0 \leq p_j \leq 1$ (e.g. implemented through a sigmoid nonlinearity as the last layer of the *DecoderNeuralNet*(.)), where D is the dimensionality of \mathbf{x} ,
- Bernoulli*(.; p) is the probability mass function (PMF) of the Bernoulli distribution.

Intractabilities

- The main difficulty of maximum likelihood learning in DLVMs is that the marginal probability of data under the model is typically **intractable**.
- This is due to the integral in equation for computing the marginal likelihood (or model evidence)
- Note that the joint distribution $p_{\theta}(\mathbf{x}, \mathbf{z})$ is efficient to compute, and that the densities are related through the basic identity:

$$p_{\theta}(\mathbf{z}|\mathbf{x}) = \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{p_{\theta}(\mathbf{x})}$$

- Since $p_{\theta}(\mathbf{x}, \mathbf{z})$ is tractable to compute
- a tractable marginal likelihood $p_{\theta}(\mathbf{x})$ leads to a tractable posterior $p_{\theta}(\mathbf{z}|\mathbf{x})$, and vice versa
- But both are intractable
- Approximate inference techniques allow us to approximate the posterior $p_{\theta}(\mathbf{z}|\mathbf{x})$ and the marginal likelihood $p_{\theta}(\mathbf{x})$ in DLVMs.

Variational Autoencoder

Encoder or approximate posterior

- Let $p_{\theta}(\mathbf{x}, \mathbf{z})$ be a latent-variable model with observed variables \mathbf{x} and latent variables \mathbf{z}
- To turn a DLVMs intractable posterior inference and learning problems into tractable problems, we introduce a parametric inference model $q_{\phi}(\mathbf{z}|\mathbf{x})$
- This model is also called an encoder.
- With ϕ we indicate the parameters of this inference model, also called the variational parameters.
- We optimize the variational parameters such that:

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \approx p_{\theta}(\mathbf{z}|\mathbf{x})$$

- As we will explain, this approximation to the posterior help us optimize the marginal likelihood.

- Like a DLVM, the inference model can be (almost) any directed graphical model:

$$q_{\phi}(\mathbf{z}|\mathbf{x}) = q_{\phi}(\mathbf{z}_1, \dots, \mathbf{z}_M|\mathbf{x}) = \prod_{j=1}^M q_{\phi}(\mathbf{z}_j | Pa(\mathbf{z}_j), \mathbf{x})$$

- similar to a DLVM, the distribution $q_{\phi}(\mathbf{z}|\mathbf{x})$ can be parameterized using deep neural networks
- In this case, the variational parameters ϕ include the weights and biases of the neural network

$$\begin{aligned}(\mu, \log \sigma) &= \text{EncoderNeuralNet}_{\phi}(\mathbf{x}) \\ q_{\phi}(\mathbf{z}|\mathbf{x}) &= \mathcal{N}(\mathbf{z}; \mu, \text{diag}(\sigma))\end{aligned}$$

- we use a single encoder neural network to perform posterior inference over all of the datapoints in our dataset

Evidence lower bound for VAE

- The **optimization objective** of the variational autoencoder, like in other variational methods, is the evidence lower bound
- Also called Variational lower bound
- For any choice of inference model $q_{\phi}(\mathbf{z}|\mathbf{x})$, including the choice of variational parameters , we have:

$$\begin{aligned}\log p_{\theta}(\mathbf{x}) &= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x})] \\&= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right] \right] \\&= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right] \right] \\&= \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right] \right]}_{=\mathcal{L}_{\theta, \phi}(\mathbf{x}) \text{ (ELBO)}} + \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\log \left[\frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \right] \right]}_{=D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x})||p_{\theta}(\mathbf{z}|\mathbf{x}))}\end{aligned}$$

- KL divergence between $q_{\phi}(\mathbf{z}|\mathbf{x})$ and $p_{\theta}(\mathbf{z}|\mathbf{x})$, which is non-negative:

$$D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x})||p_{\theta}(\mathbf{z}|\mathbf{x})) \geq 0$$


- and zero if, and only if, $q_{\phi}(\mathbf{z}|\mathbf{x})$ equals the true posterior distribution
- The first term in the variational lower bound, also called the evidence lower bound (ELBO):

$$\mathcal{L}_{\theta,\phi}(\mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})]$$

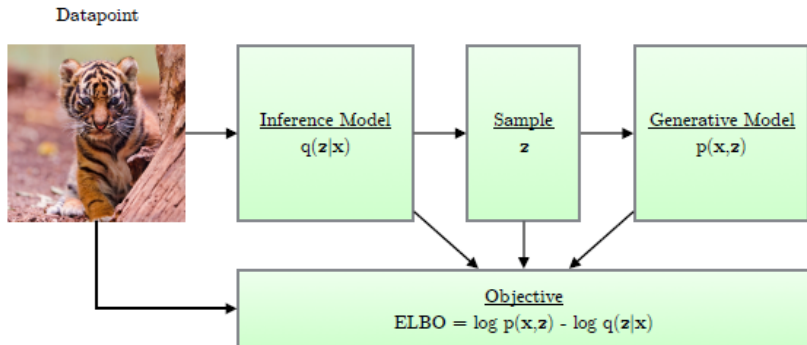
- Due to the non-negativity of the *KLDivergence*, the ELBO is a lower bound on the log-likelihood of the data.

$$\begin{aligned} \mathcal{L}_{\theta,\phi}(\mathbf{x}) &= \log p_{\theta}(\mathbf{x}) - D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x})||p_{\theta}(\mathbf{z}|\mathbf{x})) \\ &\leq \log p_{\theta}(\mathbf{x}) \end{aligned}$$

VAE in a nutshell



Control flow in VAE



A Double-Edged Sword

- Maximization of the ELBO $L_{\theta,\phi}(\mathbf{x})$ w.r.t. the parameters θ and ϕ , will concurrently optimize the two things we care about:
- It will approximately maximize the marginal likelihood $p_{\theta}(\mathbf{x})$. This means that our generative model will become better.
- It will minimize the KL divergence of the approximation $q_{\phi}(\mathbf{z}|\mathbf{x})$ from the true posterior $p_{\theta}(\mathbf{z}|\mathbf{x})$, so $q_{\phi}(\mathbf{z}|\mathbf{x})$ becomes better.

Stochastic gradient-based optimization of the ELBO

- An important property of the ELBO, is that it allows joint optimization w.r.t. all parameters (ϕ and θ) using SGD.
- We can start out with random initial values of ϕ and θ , and stochastically optimize their values until convergence.
- Given a dataset with i.i.d. data, the ELBO objective is the sum (or average) of individual-datapoint ELBOs:

$$\mathcal{L}_{\theta,\phi}(\mathcal{D}) = \sum_{x \in \mathcal{D}} \mathcal{L}_{\theta,\phi}(x)$$

- Unbiased gradients of the ELBO w.r.t. the generative model parameters are simple to obtain:

$$\begin{aligned}
 \nabla_{\theta} \mathcal{L}_{\theta, \phi}(\mathbf{x}) &= \nabla_{\theta} \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})] \\
 &= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\nabla_{\theta} (\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x}))] \\
 &\simeq \nabla_{\theta} (\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})) \\
 &= \nabla_{\theta} (\log p_{\theta}(\mathbf{x}, \mathbf{z}))
 \end{aligned}$$

- Unbiased gradients w.r.t. the variational parameters are more difficult to obtain, since the ELBOs expectation is taken w.r.t. the distribution $q_{\phi}(\mathbf{z}|\mathbf{x})$, which is a function of ϕ . I.e., in general:

$$\begin{aligned}
 \nabla_{\phi} \mathcal{L}_{\theta, \phi}(\mathbf{x}) &= \nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})] \\
 &\neq \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\nabla_{\phi} (\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x}))]
 \end{aligned}$$

Reparameterization trick

Change of variables

- For continuous latent variables and a differentiable encoder and generative model, the ELBO can be straightforwardly differentiated w.r.t. both θ and ϕ through a **change of variables**, also called the **reparameterization trick**
- First, we express the random variable \mathbf{z} as some differentiable (and invertible) transformation of another random variable ϵ , given \mathbf{z} and

$$\mathbf{z} = \mathbf{g}(\epsilon, \phi, \mathbf{x})$$

- where the distribution of random variable ϵ is independent of \mathbf{x} or ϕ

Gradient of expectation under change of variable

- Given such a change of variable, expectations can be rewritten in terms of ϵ

$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [f(\mathbf{z})] = \mathbb{E}_{p(\epsilon)} [f(\mathbf{z})]$$

- where $\mathbf{z} = \mathbf{g}(\epsilon, \phi, \mathbf{x})$ and the expectation and gradient operators become commutative, and we can form a simple Monte Carlo estimator:


$$\begin{aligned}\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [f(\mathbf{z})] &= \nabla_{\phi} \mathbb{E}_{p(\epsilon)} [f(\mathbf{z})] \\ &= \mathbb{E}_{p(\epsilon)} [\nabla_{\phi} f(\mathbf{z})] \\ &\simeq \nabla_{\phi} f(\mathbf{z})\end{aligned}$$

- where in the last line, $\mathbf{z} = \mathbf{g}(\epsilon, \phi, \mathbf{x})$ with random noise sample


$$\epsilon \sim p(\epsilon)$$

Illustration of the reparameterization trick

Original form



Reparameterized form



: Deterministic node



: Evaluation of f



: Random node



: Differentiation of f

Gradient of ELBO

- Under the reparameterization, we can replace an expectation w.r.t. $q_{\phi}(\mathbf{z}|\mathbf{x})$ with one w.r.t. $p(\epsilon)$. The ELBO can be rewritten as:

$$\begin{aligned}\mathcal{L}_{\theta,\phi}(\mathbf{x}) &= \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_{p(\epsilon)} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})]\end{aligned}$$

where $\mathbf{z} = g(\epsilon, \phi, \mathbf{x})$.

- As a result we can form a simple Monte Carlo estimator of the individual-datapoint ELBO where we use a single noise sample ϵ from $p(\epsilon)$:

$$\begin{aligned}\epsilon &\sim p(\epsilon) \\ \mathbf{z} &= g(\phi, \mathbf{x}, \epsilon) \\ \tilde{\mathcal{L}}_{\theta,\phi}(\mathbf{x}) &= \log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z}|\mathbf{x})\end{aligned}$$

Auto-Encoding Variational Bayes (AEVB) algorithm

Algorithm 1: Stochastic optimization of the ELBO. Since noise originates from both the minibatch sampling and sampling of $p(\epsilon)$, this is a doubly stochastic optimization procedure. We also refer to this procedure as the *Auto-Encoding Variational Bayes* (AEVB) algorithm.

Data:

\mathcal{D} : Dataset

$q_{\phi}(z|x)$: Inference model

$p_{\theta}(x, z)$: Generative model

Result:

θ, ϕ : Learned parameters

$(\theta, \phi) \leftarrow$ Initialize parameters

while *SGD not converged* **do**

$\mathcal{M} \sim \mathcal{D}$ (Random minibatch of data)

$\epsilon \sim p(\epsilon)$ (Random noise for every datapoint in \mathcal{M})

 Compute $\tilde{\mathcal{L}}_{\theta, \phi}(\mathcal{M}, \epsilon)$ and its gradients $\nabla_{\theta, \phi} \tilde{\mathcal{L}}_{\theta, \phi}(\mathcal{M}, \epsilon)$

 Update θ and ϕ using SGD optimizer

end

Computation of $\log q_{\phi}(\mathbf{z}|\mathbf{x})$

- Note that we typically know the density $p(\epsilon)$, since this is the density of the chosen noise distribution. As long as $\mathbf{g}(\cdot)$ is an invertible function, the densities of ϵ and \mathbf{z} are related as:

$$\log q_{\phi}(\mathbf{z}|\mathbf{x}) = \log p(\epsilon) - \log d_{\phi}(\mathbf{x}, \epsilon)$$

where

$$\log d_{\phi}(\mathbf{x}, \epsilon) = \log \left| \det \left(\frac{\partial \mathbf{z}}{\partial \epsilon} \right) \right|$$

- The Jacobian matrix contains all first derivatives of the transformation from ϵ to \mathbf{z} :

$$\frac{\partial \mathbf{z}}{\partial \epsilon} = \frac{\partial (z_1, \dots, z_k)}{\partial (\epsilon_1, \dots, \epsilon_k)} = \begin{pmatrix} \frac{\partial z_1}{\partial \epsilon_1} & \dots & \frac{\partial z_1}{\partial \epsilon_k} \\ \vdots & \ddots & \vdots \\ \frac{\partial z_k}{\partial \epsilon_1} & \dots & \frac{\partial z_k}{\partial \epsilon_k} \end{pmatrix}$$

- we can build very expressive transformations $g()$ for which $\log d_\phi(\mathbf{x}, \boldsymbol{\epsilon})$ is simple to compute, resulting in highly expressive inference models $q_\phi(\mathbf{z}|\mathbf{x})$.

Factorized gaussian posteriors

A common choice is a simple factorized Gaussian encoder $q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}, \text{diag}(\boldsymbol{\sigma}^2))$:

$$(\boldsymbol{\mu}, \log \boldsymbol{\sigma}) = \text{EncoderNeuralNet}_{\phi}(\mathbf{x})$$

$$q_{\phi}(\mathbf{z}|\mathbf{x}) = \prod_i q_{\phi}(z_i|\mathbf{x}) = \prod_i \mathcal{N}(z_i; \mu_i, \sigma_i^2)$$

where $\mathcal{N}(z_i; \mu_i, \sigma_i^2)$ is the PDF of the univariate Gaussian distribution.

- After reparameterization, we can write:

$$\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$(\boldsymbol{\mu}, \log \boldsymbol{\sigma}) = \text{EncoderNeuralNet}_{\phi}(\mathbf{x})$$

$$\mathbf{z} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon}$$

- The Jacobian of the transformation from ϵ to \mathbf{z} is:

$$\frac{\partial \mathbf{z}}{\partial \epsilon} = \text{diag}(\sigma),$$

i.e. a diagonal matrix with the elements of $\text{diag}(\sigma)$ on the diagonal. The determinant of a diagonal (or more generally, triangular) matrix is the product of its diagonal terms. The log determinant of the Jacobian is therefore:

$$\log d_{\phi}(\mathbf{x}, \epsilon) = \log \left| \det \left(\frac{\partial \mathbf{z}}{\partial \epsilon} \right) \right| = \sum_i \log \sigma_i$$

and the posterior density is:

$$\begin{aligned} \log q_{\phi}(\mathbf{z}|\mathbf{x}) &= \log p(\epsilon) - \log d_{\phi}(\mathbf{x}, \epsilon) \\ &= \sum_i \log \mathcal{N}(\epsilon_i; 0, 1) - \log \sigma_i \end{aligned}$$

when $\mathbf{z} = \mathbf{g}(\epsilon, \phi, \mathbf{x})$.

Full-covariance Gaussian posterior

- The factorized Gaussian posterior can be extended to a Gaussian with full covariance:

$$q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$$

A reparameterization of this distribution is given by:

$$\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$\mathbf{z} = \boldsymbol{\mu} + \mathbf{L}\boldsymbol{\epsilon}$$

- where \mathbf{L} is a lower (or upper) triangular matrix, with non-zero entries on the diagonal.

The reason for this parameterization of the full-covariance Gaussian, is that the Jacobian determinant is remarkably simple. The Jacobian in this case is trivial: $\frac{\partial \mathbf{z}}{\partial \boldsymbol{\epsilon}} = \mathbf{L}$. Note that the determinant of a triangular matrix is the product of its diagonal elements. Therefore, in this parameterization:

$$\log \left| \det \left(\frac{\partial \mathbf{z}}{\partial \boldsymbol{\epsilon}} \right) \right| = \sum_i \log |L_{ii}|$$

And the log-density of the posterior is:

$$\log q_{\phi}(\mathbf{z}|\mathbf{x}) = \log p(\boldsymbol{\epsilon}) - \sum_i \log |L_{ii}|$$

This parameterization corresponds to the Cholesky decomposition $\boldsymbol{\Sigma} = \mathbf{L}\mathbf{L}^T$ of the covariance of \mathbf{z} :

$$\begin{aligned}\boldsymbol{\Sigma} &= \mathbb{E} \left[(\mathbf{z} - \mathbb{E}[\mathbf{z}]) (\mathbf{z} - \mathbb{E}[\mathbf{z}])^T \right] \\ &= \mathbb{E} \left[\mathbf{L}\boldsymbol{\epsilon} (\mathbf{L}\boldsymbol{\epsilon})^T \right] = \mathbf{L} \mathbb{E} \left[\boldsymbol{\epsilon} \boldsymbol{\epsilon}^T \right] \mathbf{L}^T \\ &= \mathbf{L}\mathbf{L}^T\end{aligned}$$

Note that $\mathbb{E} [\boldsymbol{\epsilon} \boldsymbol{\epsilon}^T] = \mathbf{I}$ since $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.

One way to build a matrix \mathbf{L} with the desired properties, namely triangularity and non-zero diagonal entries, is by constructing it as follows:

$$(\mu, \log \sigma, \mathbf{L}') \leftarrow \text{EncoderNeuralNet}_{\phi}(\mathbf{x})$$

$$\mathbf{L} \leftarrow \mathbf{L}_{mask} \odot \mathbf{L}' + \text{diag}(\sigma)$$

\mathbf{L}_{mask} is a masking matrix with zeros on and above the diagonal, and ones below the diagonal. The log-determinant is identical to the factorized Gaussian case:

$$\log \left| \det \left(\frac{\partial \mathbf{z}}{\partial \epsilon} \right) \right| = \sum_i \log \sigma_i$$

VAE with a full-covariance Gaussian inference model

Algorithm 2: Computation of unbiased estimate of single-datapoint ELBO for example VAE with a full-covariance Gaussian inference model and a factorized Bernoulli generative model. L_{mask} is a masking matrix with zeros on and above the diagonal, and ones below the diagonal. Note that L must be a triangular matrix with positive entries on the diagonal.

Data:

x : a datapoint, and optionally other conditioning information

ϵ : a random sample from $p(\epsilon) = \mathcal{N}(0, I)$

θ : Generative model parameters

ϕ : Inference model parameters

$q_\phi(z|x)$: Inference model

$p_\theta(x, z)$: Generative model

Result:

$\hat{\mathcal{L}}$: unbiased estimate of the single-datapoint ELBO $\mathcal{L}_{\theta, \phi}(x)$

$(\mu, \log \sigma, L') \leftarrow \text{EncoderNeuralNet}_\phi(x)$

$L \leftarrow L_{mask} \odot L' + \text{diag}(\sigma)$

$\epsilon \sim \mathcal{N}(0, I)$

$z \leftarrow L\epsilon + \mu$

$\log qz \leftarrow -\text{sum}(\frac{1}{2}(\epsilon^2 + \log(2\pi) + \log \sigma))$ $\triangleright = q_\phi(z|x)$

$\log pz \leftarrow -\text{sum}(\frac{1}{2}(z^2 + \log(2\pi)))$ $\triangleright = p_\theta(z)$

$p \leftarrow \text{DecoderNeuralNet}_\theta(z)$


$\log px \leftarrow \text{sum}(x \odot \log p + (1 - x) \odot \log(1 - p))$ $\triangleright = p_\theta(x|z)$

$\hat{\mathcal{L}} = \log px + \log pz - \log qz$

Estimation of the marginal likelihood

- After training a VAE, we can estimate the probability of data under the model using an importance sampling technique **Rezende et al. [2014]**

Marginal likelihood and ELBO as KL divergences



Marginal likelihood and ELBO as KL divergences

- One additional perspective is that the ELBO can be viewed as a maximum likelihood objective in an augmented space.
- For some fixed choice of encoder $q_\phi(\mathbf{z}|\mathbf{x})$, we can view the joint distribution $p_\theta(\mathbf{x}, \mathbf{z})$ as an augmented empirical distribution over the original data \mathbf{x} and (stochastic) auxiliary features \mathbf{z} associated with each datapoint.
- The model $p_\theta(\mathbf{x}, \mathbf{z})$ then defines a joint model over the original data, and the auxiliary features.

Other Examples of Variational treatment

- Variational Mixture of Gaussians
- Variational Linear Regression
- Variational Logistic Regression etc.

Bishop et al. [2006]

Thanks you.
Questions?

References

- Kingma, D.P. Variational Inference and Deep Learning: A New Synthesis (Ph.D. Thesis) (2017).
- Christopher M Bishop. Pattern recognition and machine learning, (2006).